

# Getting the Most out of HTC with Workflows

Peter van Heusden [pvh@sanbi.ac.za](mailto:pvh@sanbi.ac.za)  
South African National Bioinformatics Institute  
adapted from slides of  
Christina Koch [ckoch5@wisc.edu](mailto:ckoch5@wisc.edu)

# Why are we here?

---

# Why are we here?

---

## To do SCIENCE!!!

- A lot of science is best-done with computing – sometimes, LOTS of computing
- Science needs to be reproducible & sustainable
- And, we'd really like science to happen **fast(er)**



---

# **GETTING THE MOST OUT OF COMPUTING (FOR RESEARCH)**

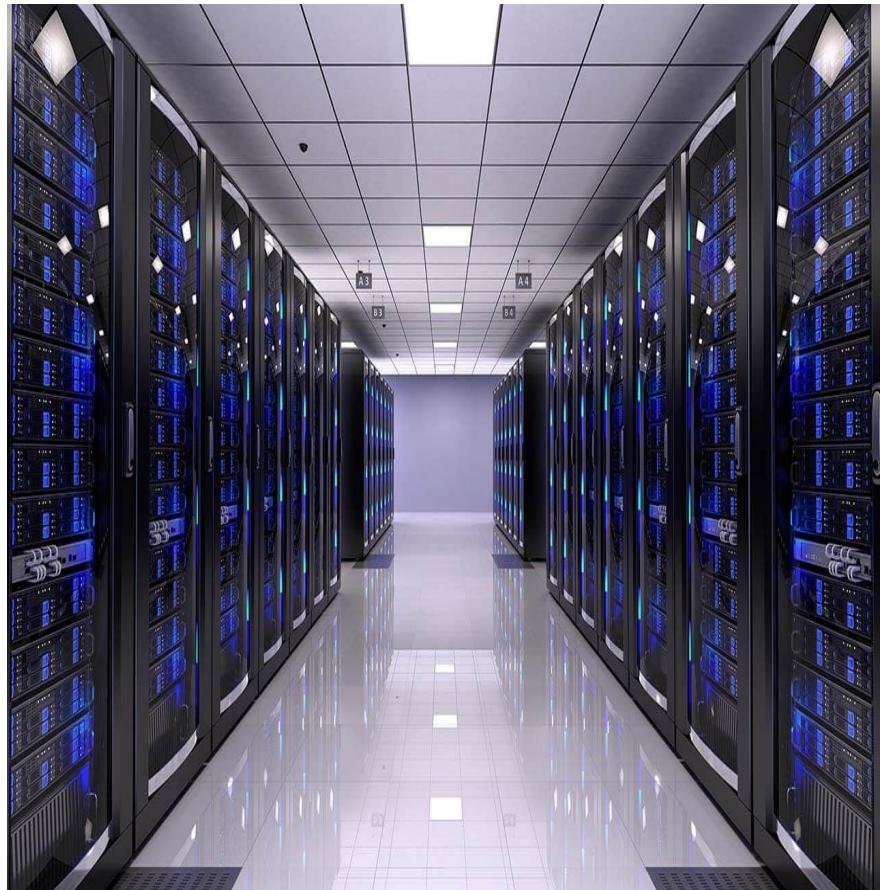
# Scaling up computing

---

From

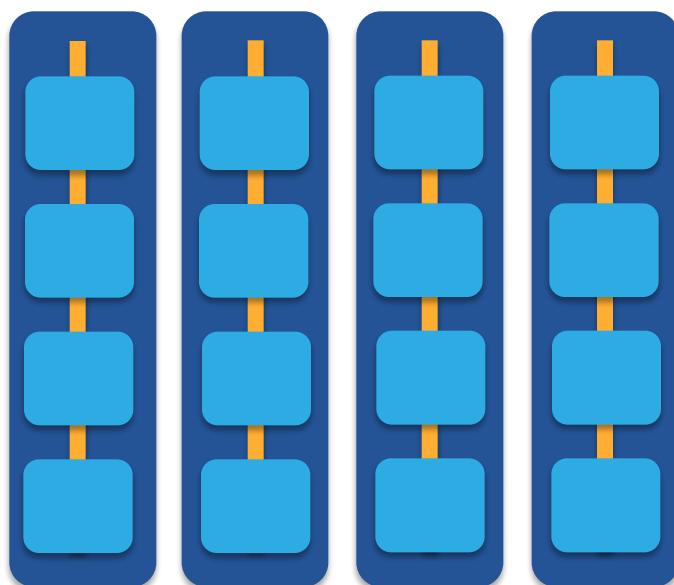


TO

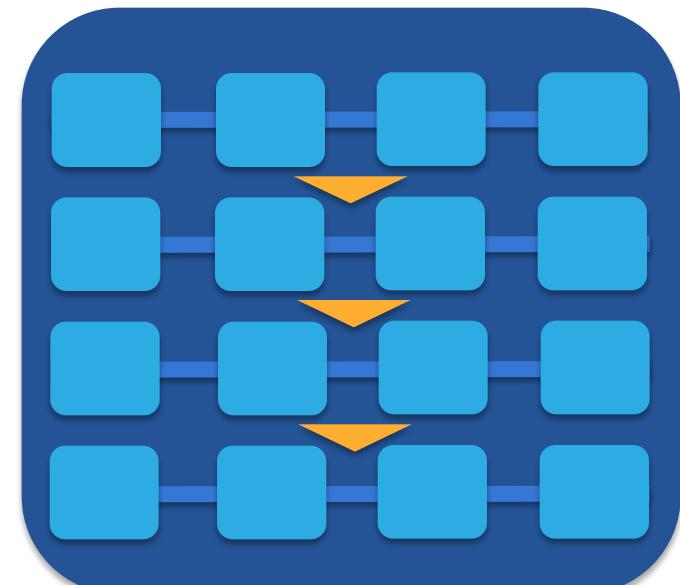


# Computing types

- Our challenge: how to make use of computers working together to tackle large compute tasks...



**high-throughput**



**high-performance (e.g. MPI)**

# Two Strategies

---

## Cloud

Focus: Service *many user groups* by providing *generic computing*

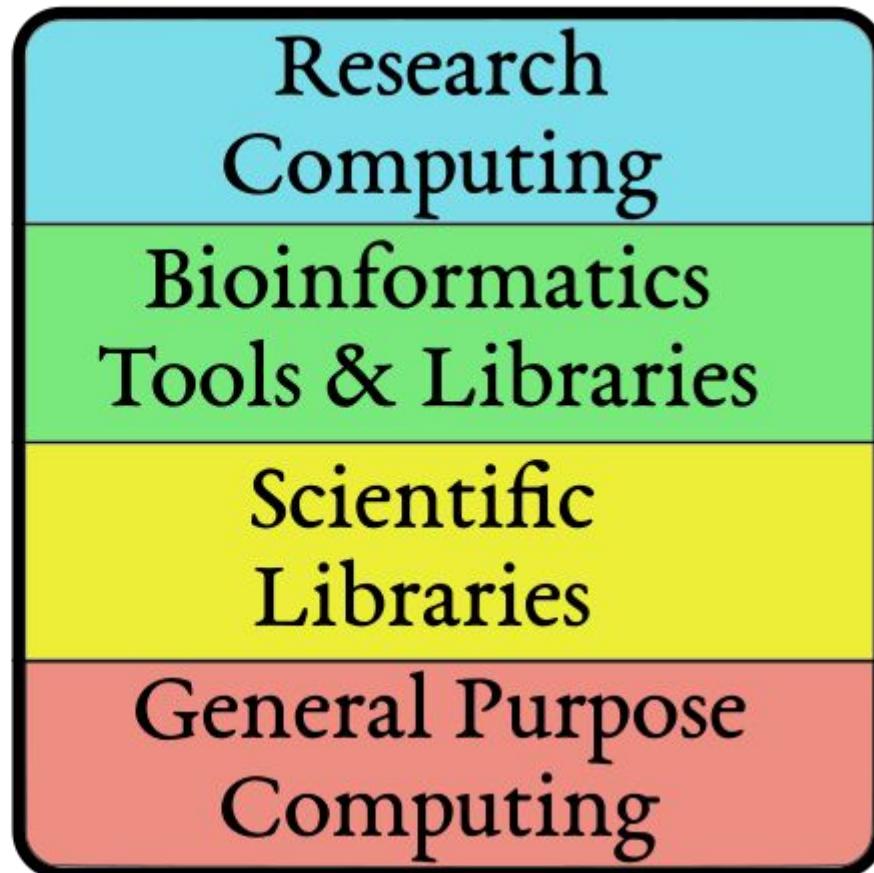
Skills focus: *systems engineering* to create *virtual infrastructure*, compose *multiple component services*

## HPC Cluster

- Focus: Service *specialised computing groups* working on *computationally challenging problems*
- Skills focus: *research software engineering* and *parallel algorithms*

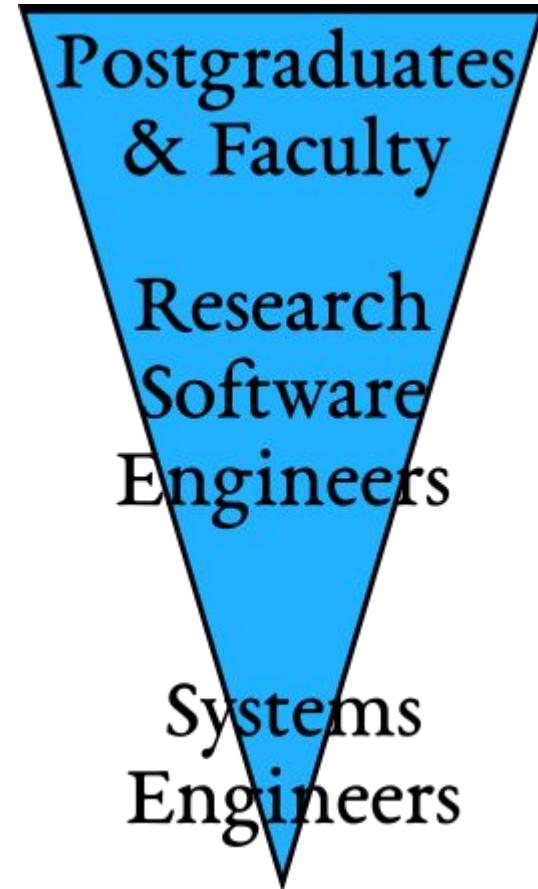
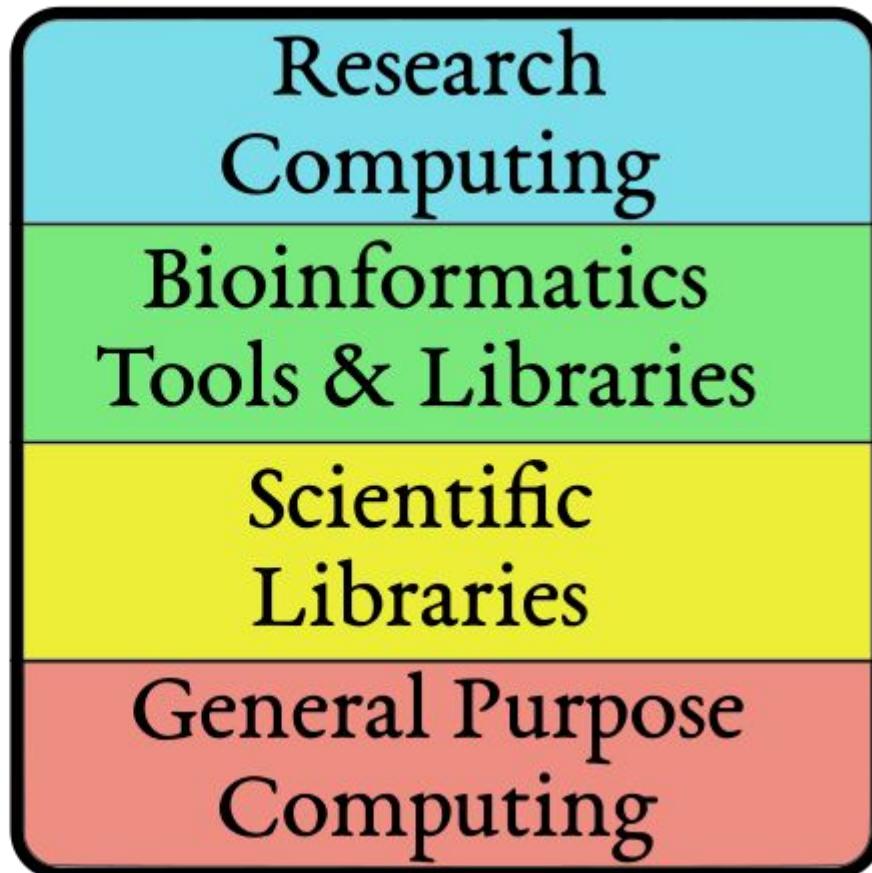
# The Research Computing Stack

---



# Research Computing Roles

---



# Two Architectures

---

## High Throughput

Focus: Workflows with many *small, largely independent* compute tasks

Optimize: *throughput*, or time from *submission* to *overall completion*

## High Performance

- Focus: Workflows with *large, highly coupled* tasks
- Optimize: *individual tasks*, software, communication between processes

# Making Good Choices

---

- How do you choose the best approach?
- Guiding question:

Is your problem “HTC-able”?



# Typical HTC Problems

---

- batches of similar program runs (>10)
- “loops” over independent tasks
- others you might not think of ...
  - programs/functions that
    - process files that are already separate
    - process columns or rows, separately
    - iterate over a parameter space
  - *a lot* of programs/functions that use multiple CPUs on the same server

**Ultimately: Can you break it up?**

# What is not HTC?

---

- fewer numbers of jobs
- jobs individually requiring significant resources
  - RAM, Data/Disk, # CPUs, time  
(though, “significant” depends on the HTC compute system you use)
- restrictive licensing

# The Real World

---

- However, it's not just about finding the right computing approach to your problem.
- These approaches will be **most** effective if they're running on appropriate compute systems.



# The Real World

---

- Not all compute systems are created equal.
- Two questions to ask:  
**What resources are available to me?**  
**Which one is the best match for the kind of computing I want to do?**

# Campus Resources

---

- Start with your local campus compute system
- Some considerations:
  - Who has access? Are there allocations?
  - What kind of system? What is it optimized for?
- An HPC cluster may not handle lots of jobs well, in the same way that an HTC system has limited multicore capabilities - be aware of how a system matches/doesn't match your computation strategy.
- Ask questions! Be a good citizen!
- If local resources are limited, explore other options.

# Beyond your campus

---

- Hosted Resources



- Clouds and clusters
  - Commercial cloud systems
  - Research clouds and clusters



# The payoff

- HTC is, beyond everything, scalable
  - If you can run 10 jobs, you can run 10,000, maybe even 10 million
- Worth pursuing the right kind of resources (if you can) for the right kind of problem.



---

# **GETTING THE MOST OUT OF HTC**

# Key HTC Tactics

---

1. Increase Overall Throughput
2. Utilize Resources Efficiently!
3. Bring Dependencies With You
4. Scale Gradually, Testing Generously
5. Automate As Many Steps As Possible

# Throughput, revisited

---

- In HTC, we optimize *throughput*: time from submission to overall completion
- Instead of making individual jobs as fast as possible, optimize how long it takes for all jobs to finish.
- Time to completion includes *engineering* time
- We do this by breaking large processes into smaller pieces and re-using components

# Breaking up is hard to do...

---

- Ideally into parallel (separate) jobs
  - reduced job requirements = more matches
  - not always easy or possible
- Strategies
  - break HTC-able steps out of a single program
  - break up loops
  - break up input
- Use self-checkpointing if jobs are too long
  - Often not supported by individual applications

# Batching (Merging) is easy

---

- A single job can
  - execute multiple independent tasks
  - execute multiple short, sequential steps
  - avoid transfer of intermediate files
- Use scripts!
  - need adequate error reporting for each “step”
  - easily handle multiple commands and arguments

# Key HTC Tactics

---

1. Increase Overall Throughput
2. **Utilize Resources Efficiently!**
3. Bring Dependencies With You
4. Scale Gradually, Testing Generously
5. Automate As Many Steps As Possible

# Know and Optimize Job Use of Resources!

---

- **CPUs** (“1” is best for matching)
  - restrict, if necessary/possible
  - software that uses all available CPUs is BAD!
- **CPU Time**

> ~5 min, < ~1 day; **Ideal: 1-2 hours**
- **RAM** (not always easily modified)
- **Disk** per-job (execute) and in-total (submit)
- **Network Bandwidth**
  - minimize transfer: filter/trim/delete, compress

# Key HTC Tactics

---

1. Increase Overall Throughput
2. Utilize Resources Efficiently!
3. **Bring Dependencies With You**
4. Scale Gradually, Testing Generously
5. Automate As Many Steps As Possible

# Bring *What* with You?

---

- Software
  - Dependency management
- Data and other input files
  - Have a Research Data Management strategy
  - Record parameters
    - do you use “reference data”?
    - including random number seeds
    - note data provenance



# Each Workflow Step Has a Wrapper

---

- Before task execution
  - transfer/prepare files and directories
  - setup/configure software environment and other dependencies
- Task execution
  - prepare complex commands and arguments
  - batch together many ‘small’ tasks
- After task execution
  - filter/combine/compress files and directories
  - check for and report on errors
  - clean up temporary files

# Software Dependency Management

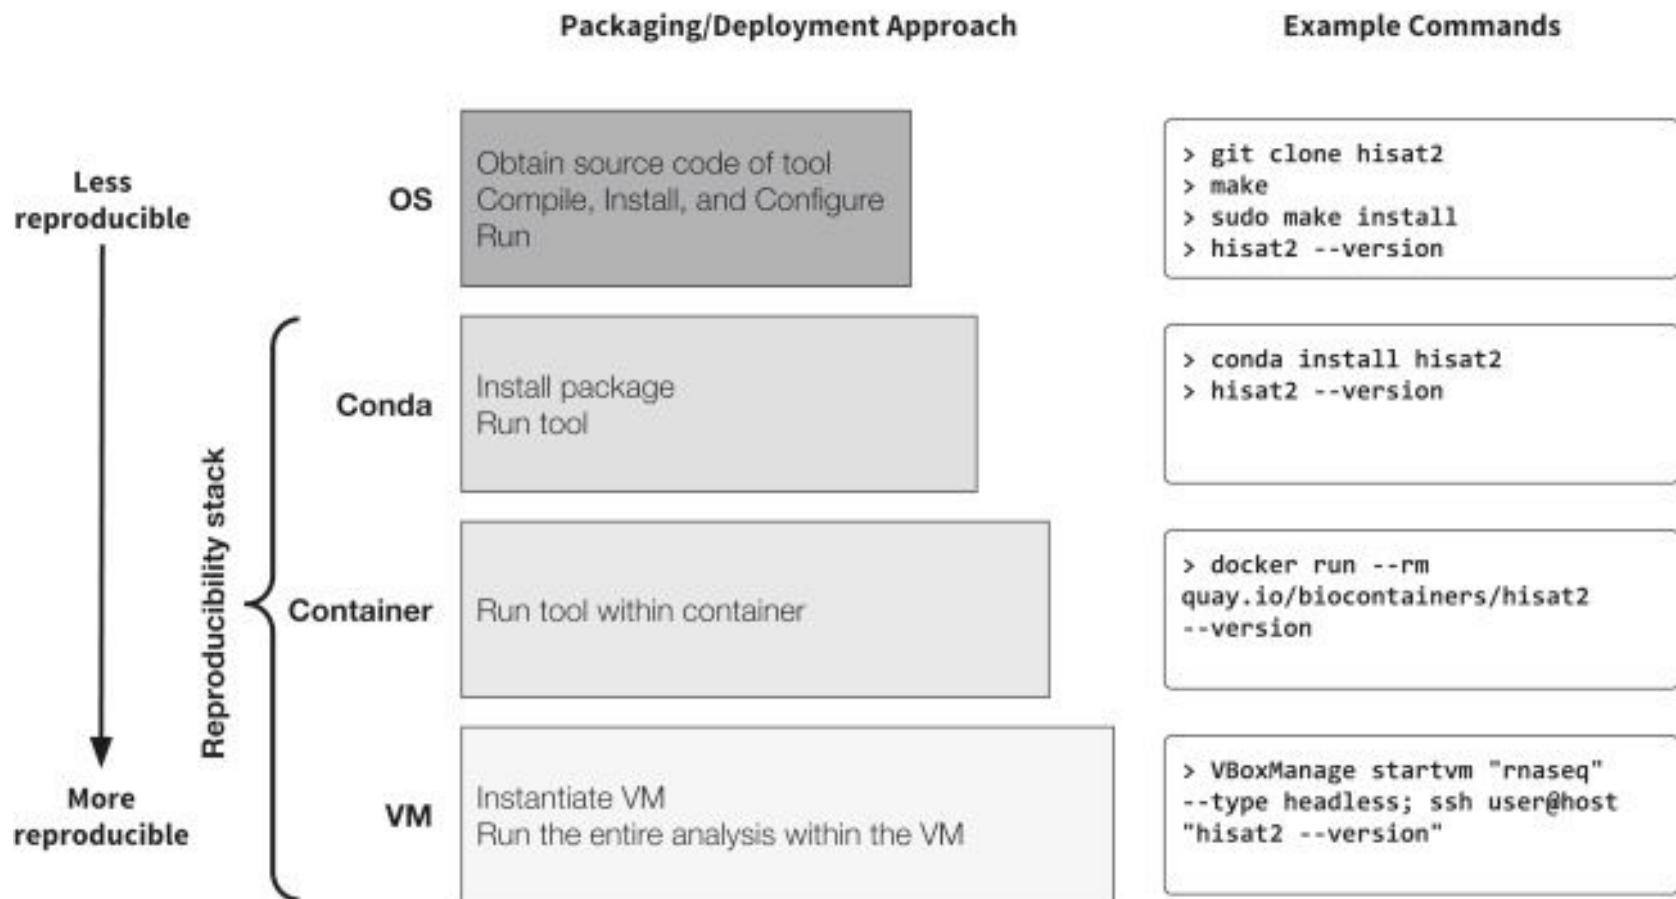


Image from Grünig et al 2018 “[Practical Computational Reproducibility in the Life Sciences](#)”

# Key HTC Tactics

---

1. Increase Overall Throughput
2. Utilize Resources Efficiently!
3. Bring Dependencies With You
- 4. Scale Gradually, Testing Generously**
5. Automate As Many Steps As Possible

# Testing, testing, testing!

---

- Allows you to optimize resource use
- Just because it worked for 10 jobs, doesn't mean it will work for 10,000 jobs (scaling issues)
  - Data transfer (in and out)
  - Discover site-specific problems

# Key HTC Tactics

---

1. Increase Overall Throughput
2. Utilize Resources Efficiently!
3. Bring Dependencies With You
4. Scale Gradually, Testing Generously
- 5. Automate As Many Steps As Possible**

# What to Automate?

---

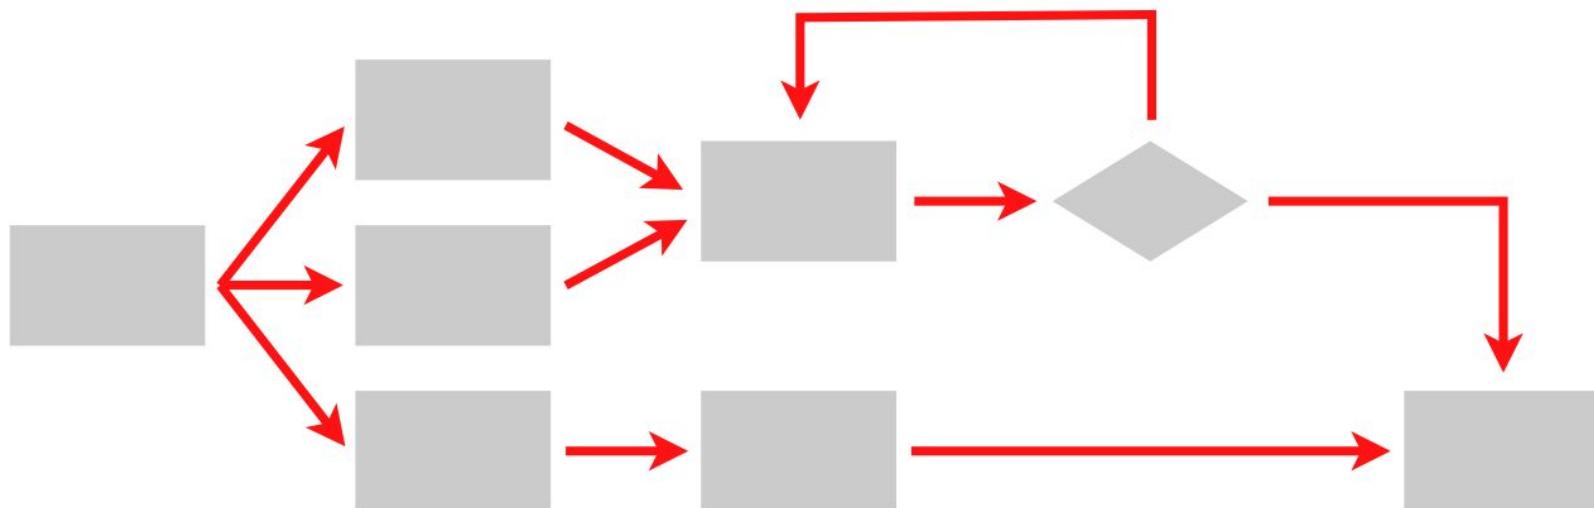
- Submitting many jobs
- Writing submit files using scripts
- Running a series of jobs, or workflow



# What is a workflow?

---

- A series of ordered steps
  - Steps
  - Connections
  - (Metadata)



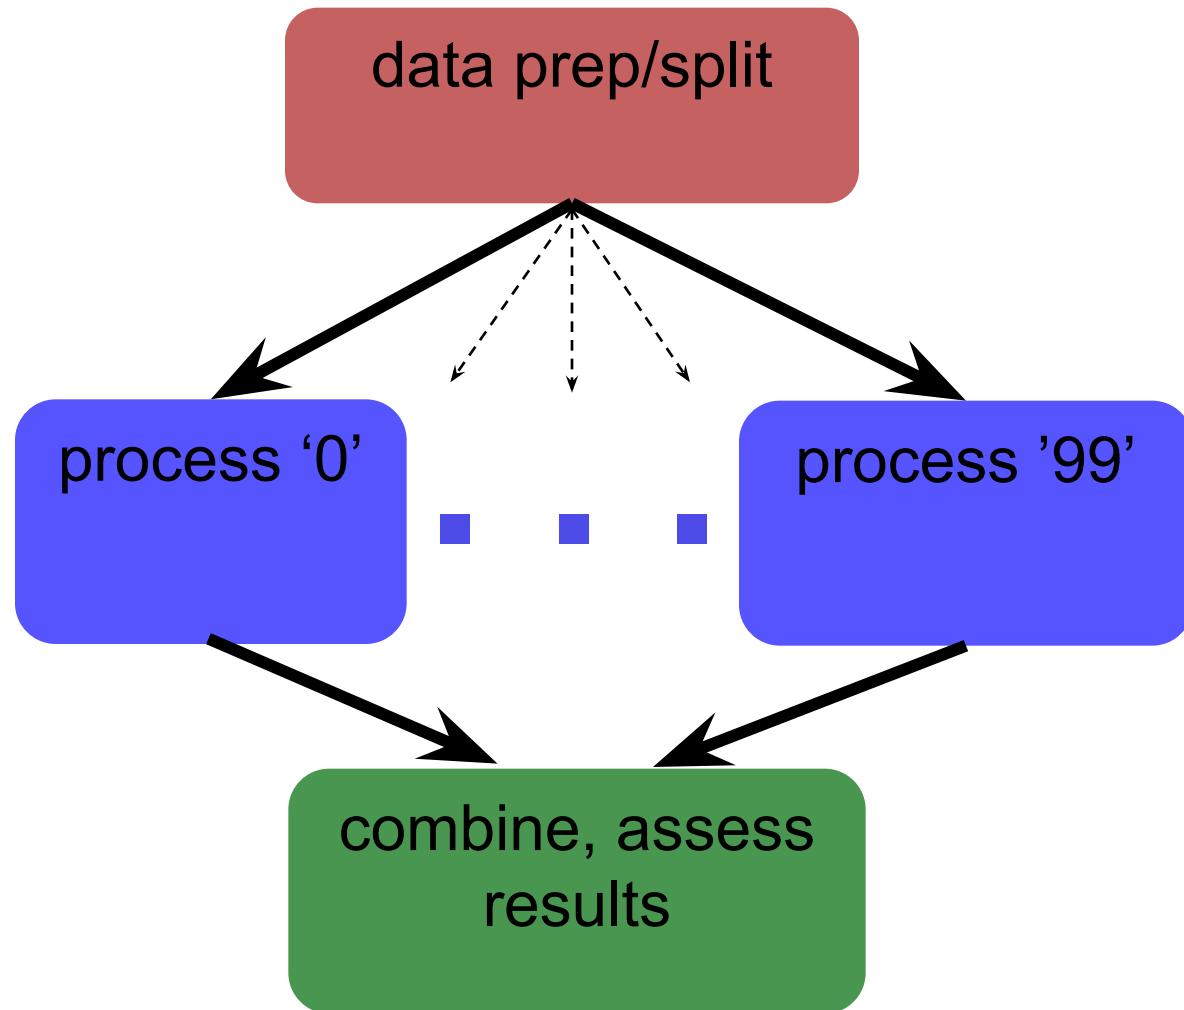
# We ❤️ workflows

- non-computing “workflows” are all around you, especially in science
  - instrument setup
  - experimental procedures and protocols
- when planned/documentated, workflows help with:
  - organizing and managing processes
  - saving time with **automation**
  - objectivity, reliability, and reproducibility  
**(THE TENETS OF GOOD SCIENCE!)**



# Scientific Workflow Management

---



# Automating workflows can save you time...

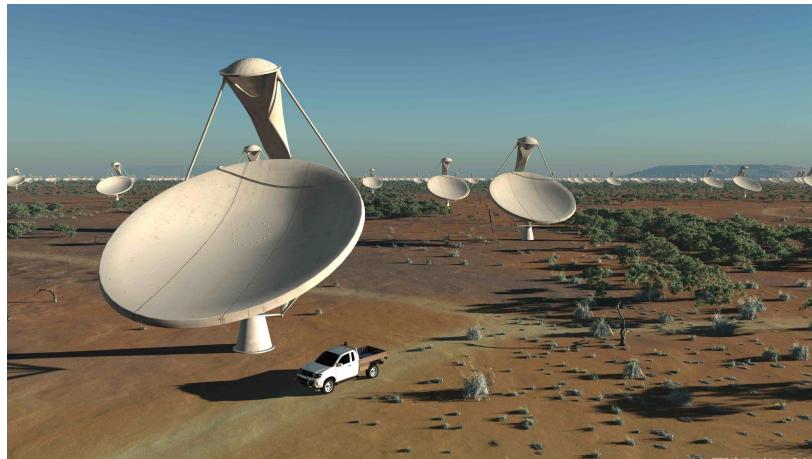
HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?  
(ACROSS FIVE YEARS)

		HOW OFTEN YOU DO THE TASK					
		50/DAY	5/DAY	DAILY	WEEKLY	MONTHLY	YEARLY
1 SECOND		1 DAY	2 HOURS	30 MINUTES	4 MINUTES	1 MINUTE	5 SECONDS
5 SECONDS		5 DAYS	12 HOURS	2 HOURS	21 MINUTES	5 MINUTES	25 SECONDS
30 SECONDS		4 WEEKS	3 DAYS	12 HOURS	2 HOURS	30 MINUTES	2 MINUTES
HOW MUCH TIME YOU SHAVE OFF	1 MINUTE	8 WEEKS	6 DAYS	1 DAY	4 HOURS	1 HOUR	5 MINUTES
	5 MINUTES	9 MONTHS	4 WEEKS	6 DAYS	21 HOURS	5 HOURS	25 MINUTES
	30 MINUTES		6 MONTHS	5 WEEKS	5 DAYS	1 DAY	2 HOURS
	1 HOUR		10 MONTHS	2 MONTHS	10 DAYS	2 DAYS	5 HOURS
	6 HOURS				2 MONTHS	2 WEEKS	1 DAY
	1 DAY					8 WEEKS	5 DAYS

# ... but there are even more benefits of automating workflows

---

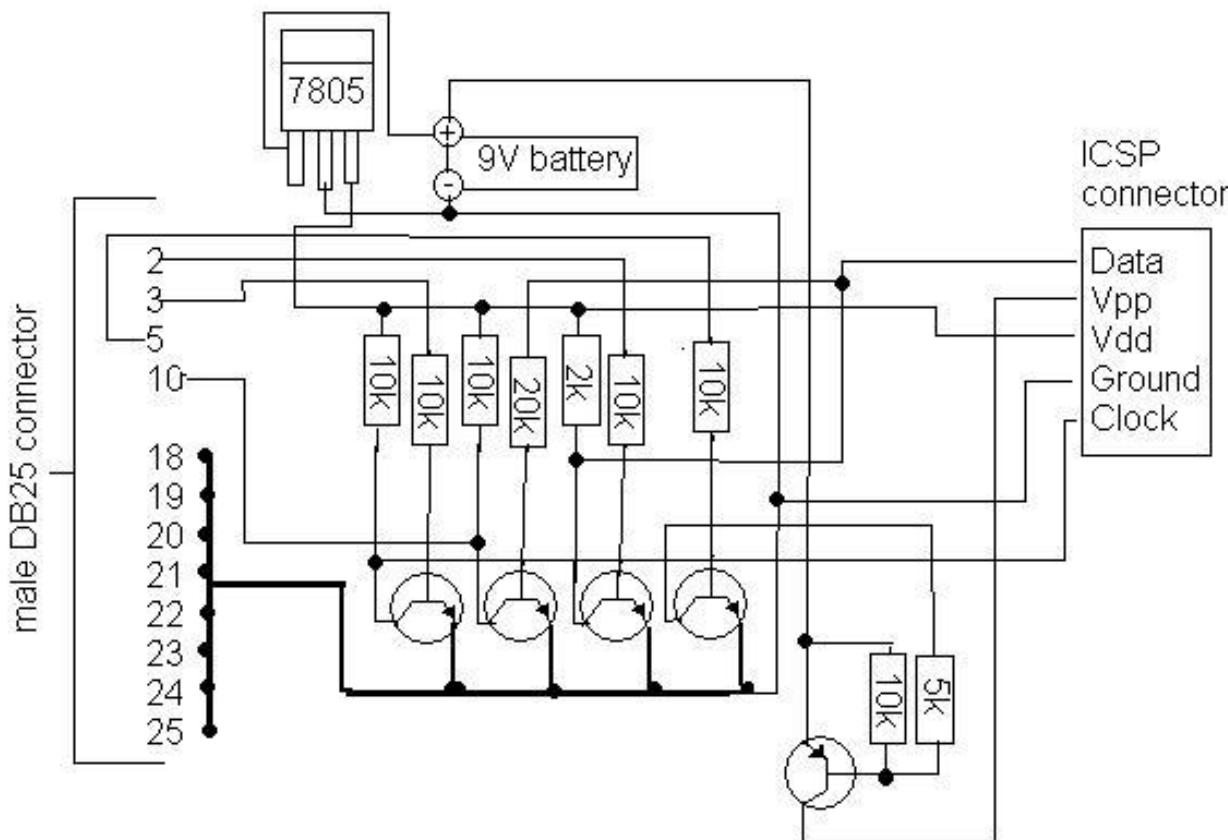
- Reproducibility
- Building knowledge and experience
- New ability to imagine greater scale, functionality, possibilities, and better SCIENCE!!



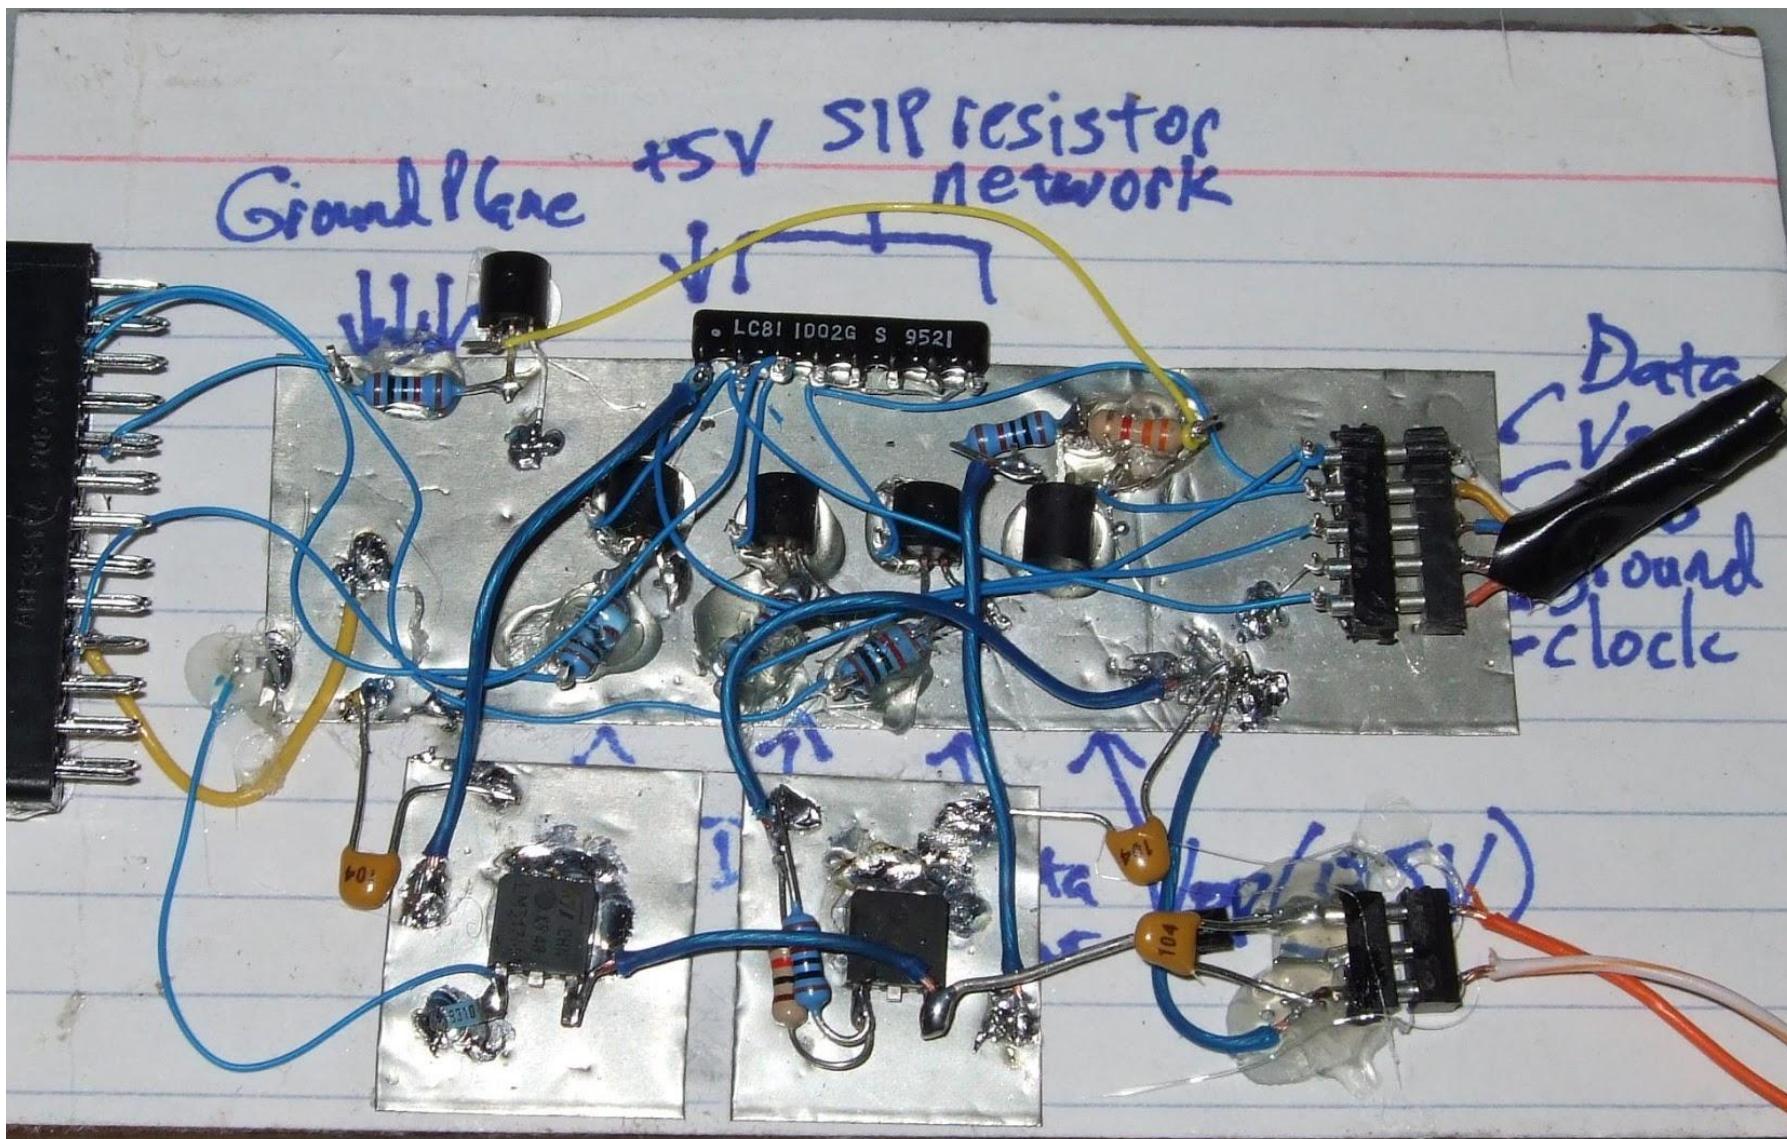
---

# **GETTING THE MOST OUT OF WORKFLOWS, PART 1**

# From schematics...



# ... to the real world



# Building a Good Workflow

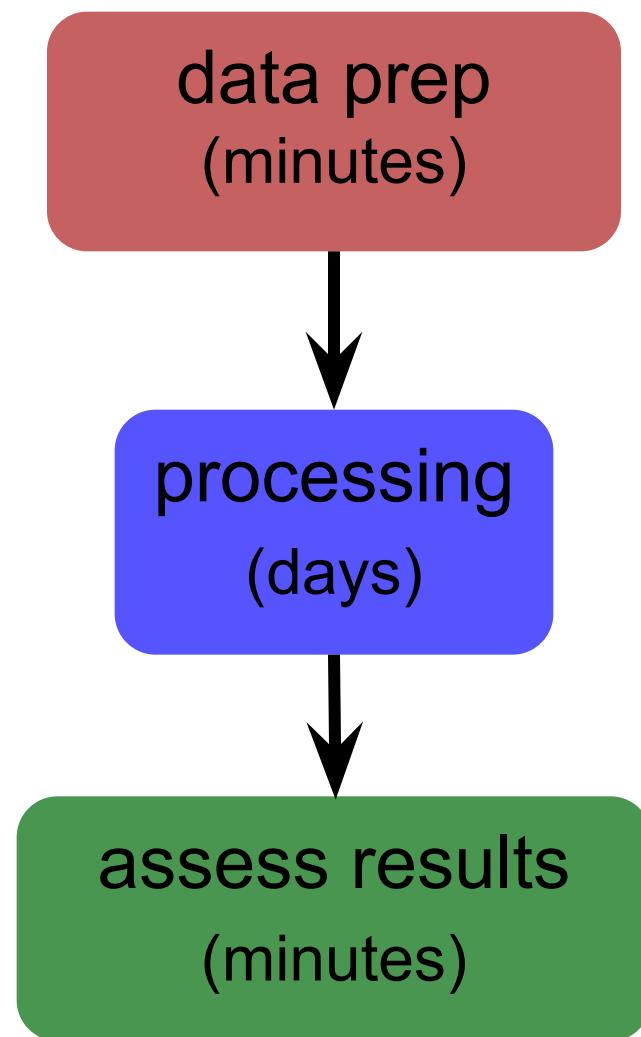
---

1. Draw out the *general workflow*
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. Make it work consistently
6. What more can you automate or error-check?
7. Publish, share and re-use

(And remember to document!)

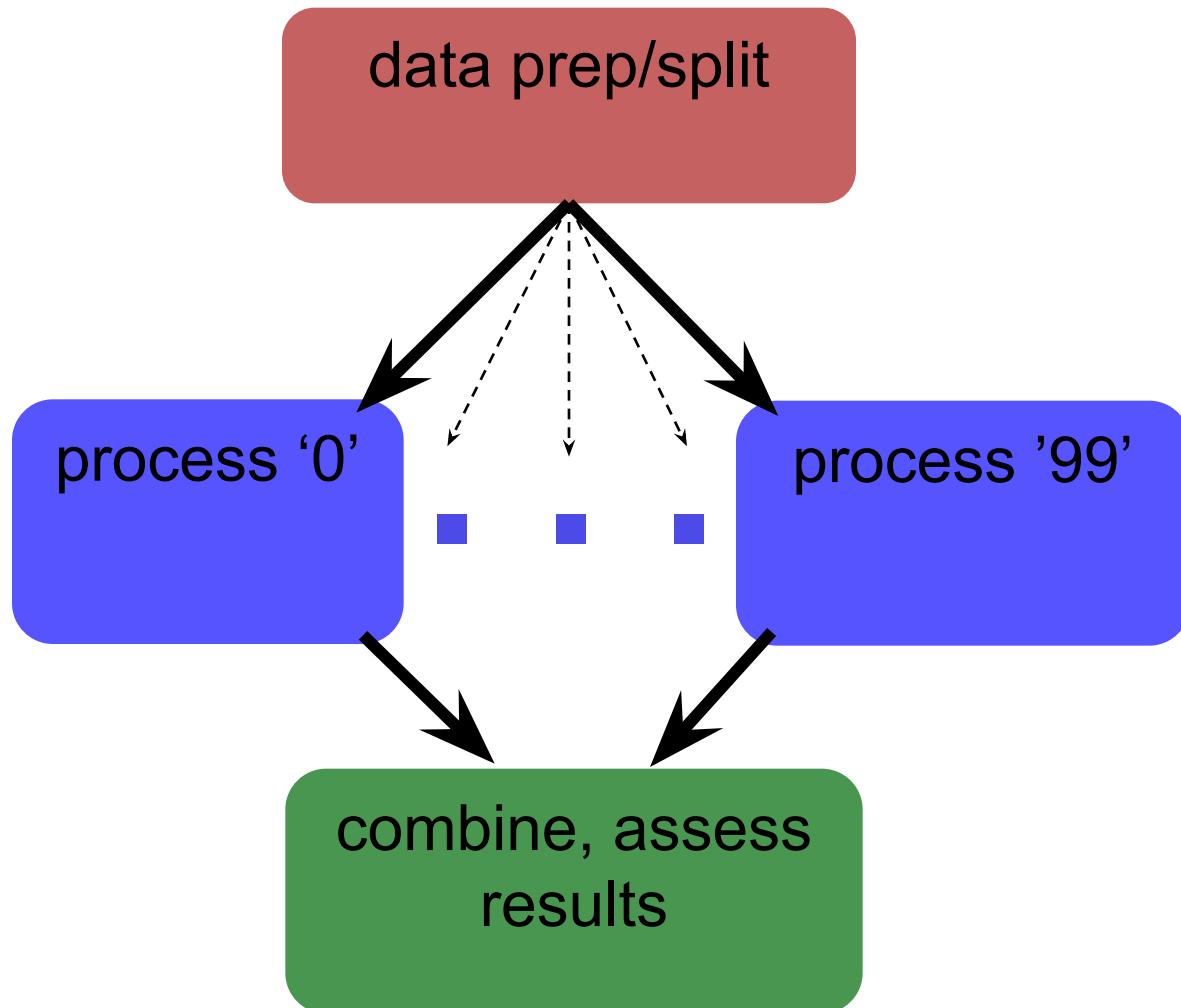
# Workflow, version 1

---



# Workflow, version 2 (HTC)

---



# Building a Good Workflow

---

1. Draw out the *general* workflow
2. **Define details (test ‘pieces’: steps / scripts)**
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. Make it work consistently
6. What more can you automate or error-check?

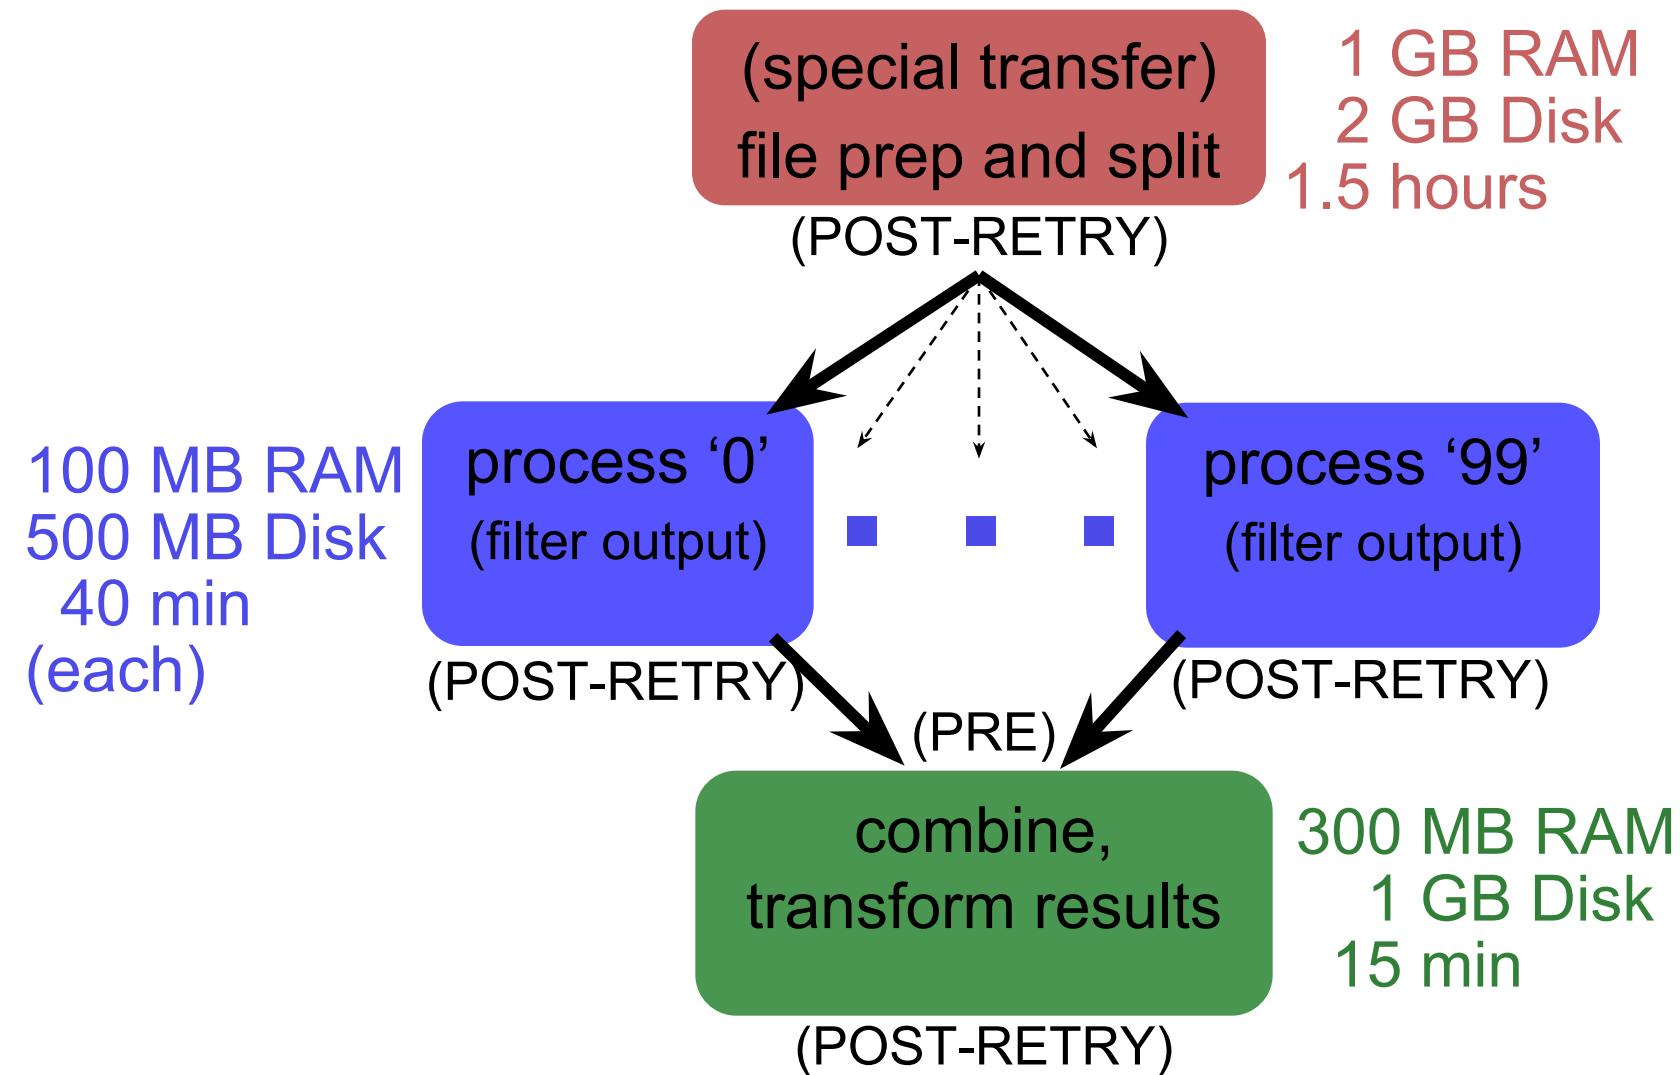
(And remember to document!)

# Determine Resource Usage

---

- Run locally first
- Then get one job running remotely
  - (on execute machine, not submit machine)!
  - get the logistics correct! (job submission, file and software setup, etc.)
- Once working, run a couple of times
  - If big variance in resource needs, should you take the...
    - Average? Median? Worst case?

# End Up with This



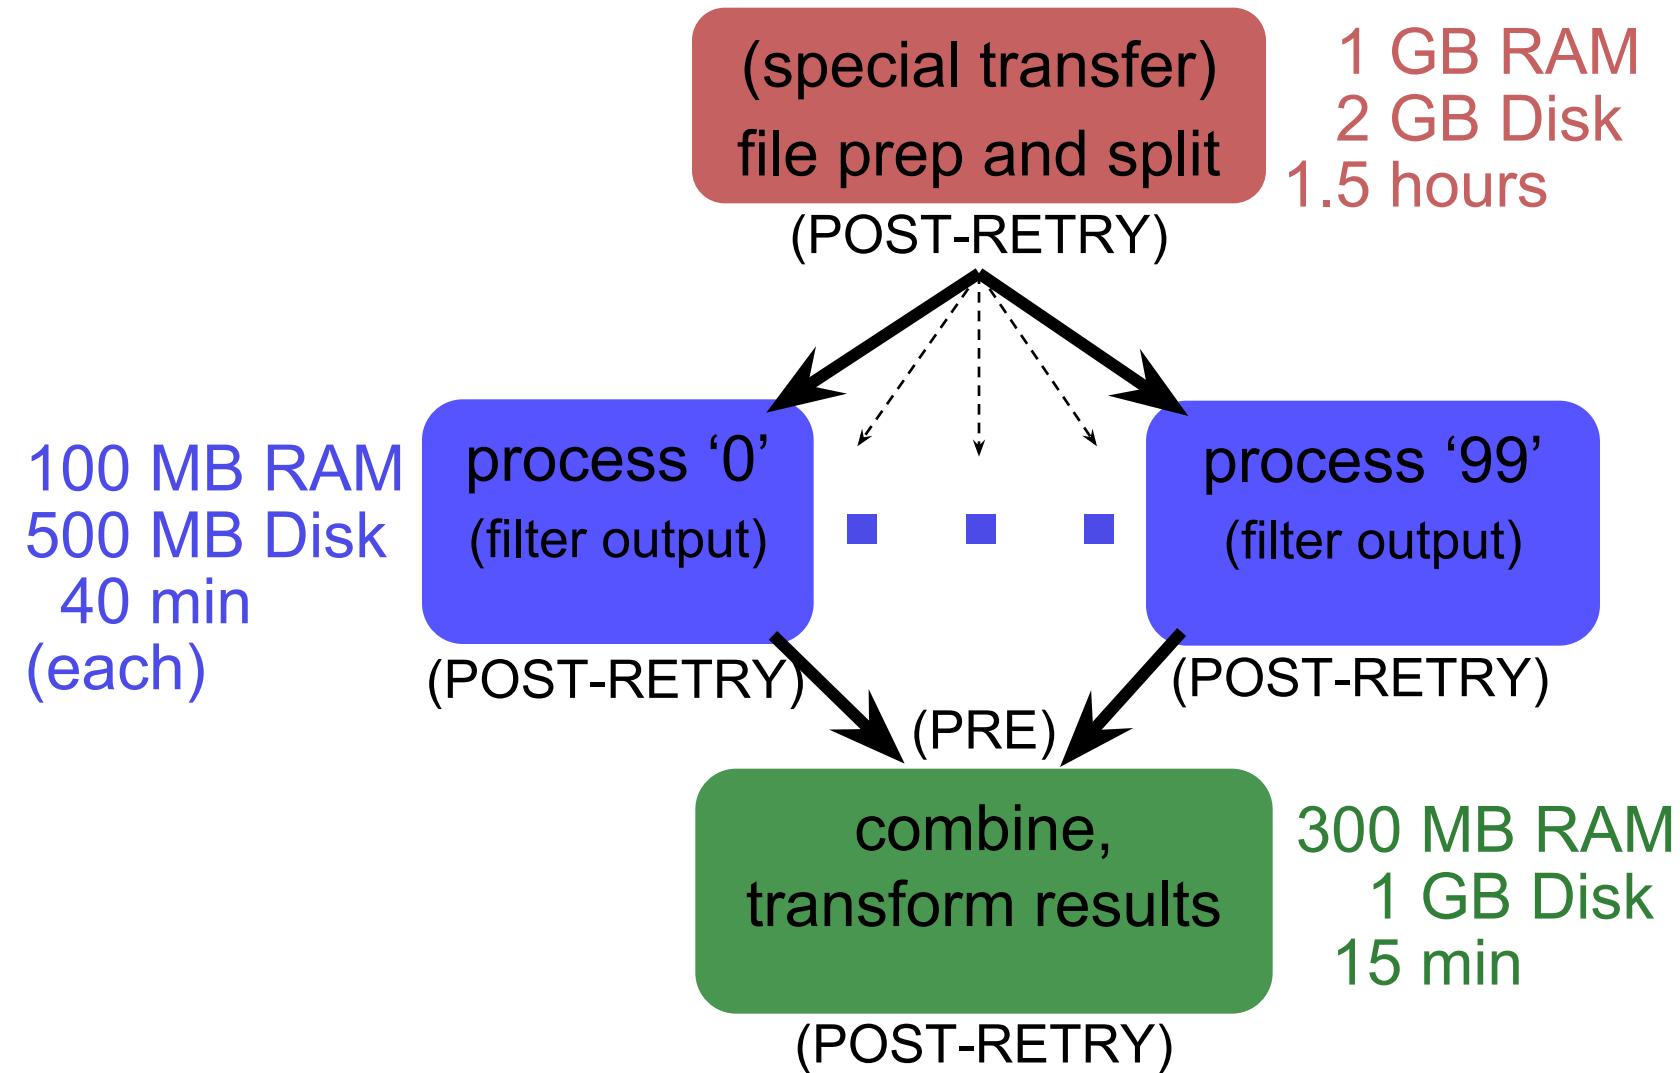
# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. **Build it modularly; test and optimize**
4. Scale-up gradually
5. Make it work consistently
6. What more can you automate or error-check?

(And remember to document!)

# To Get Here ...



# Start Here

---

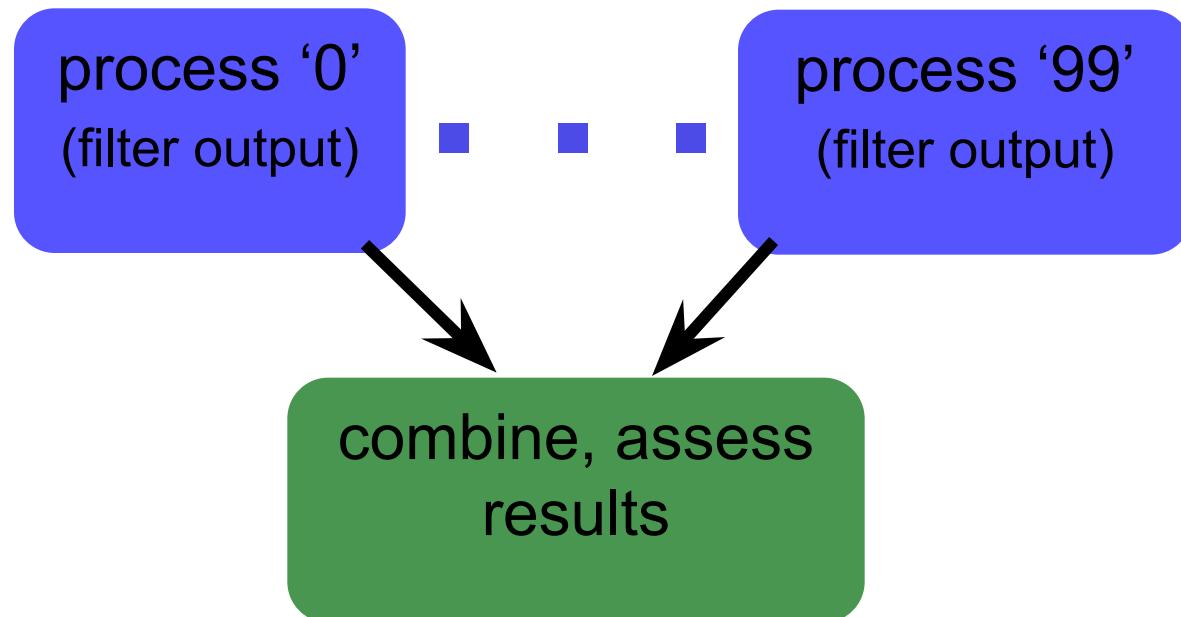
process '0'  
(filter output)



process '99'  
(filter output)

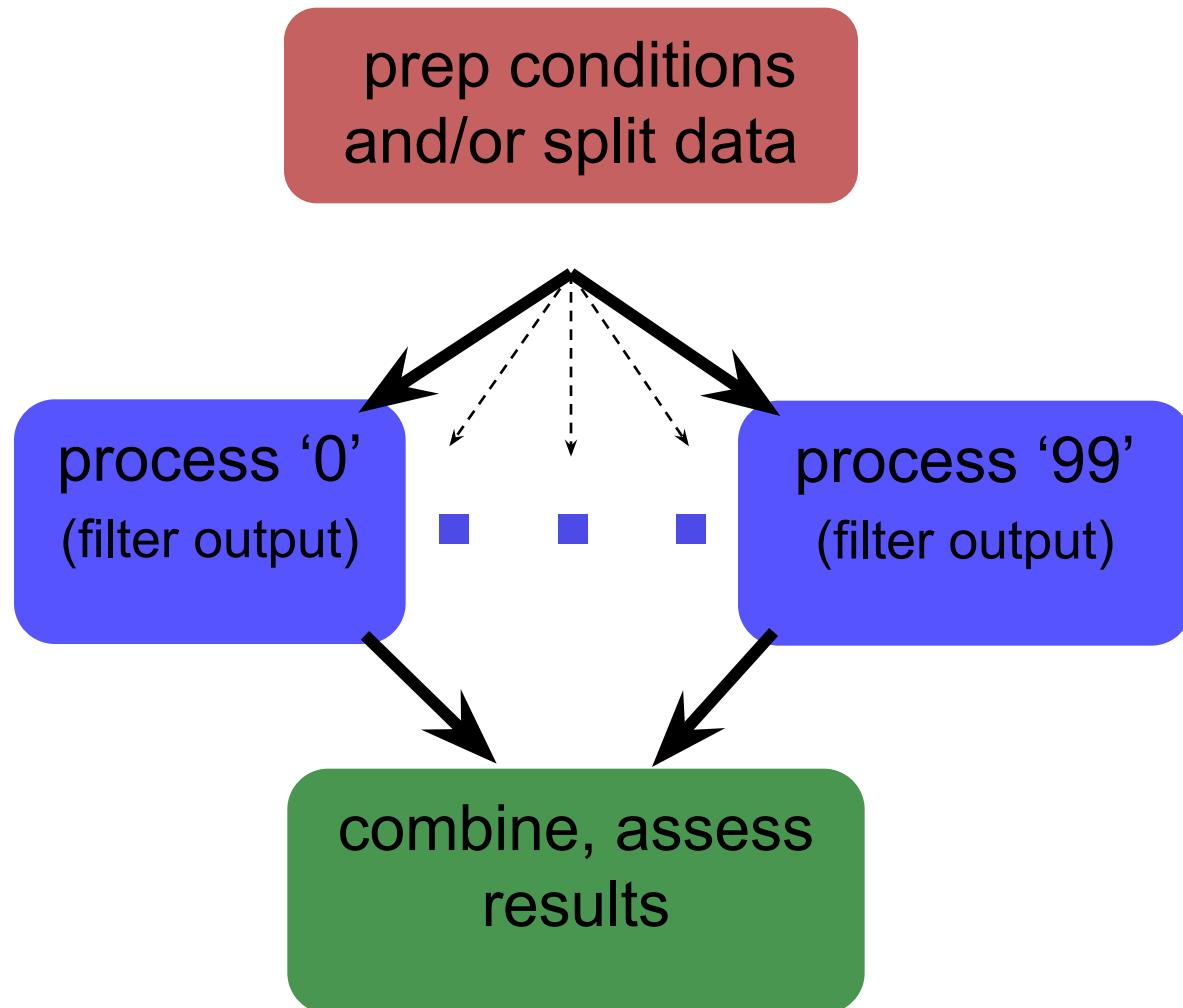
# Add a Step

---



# And Another Step

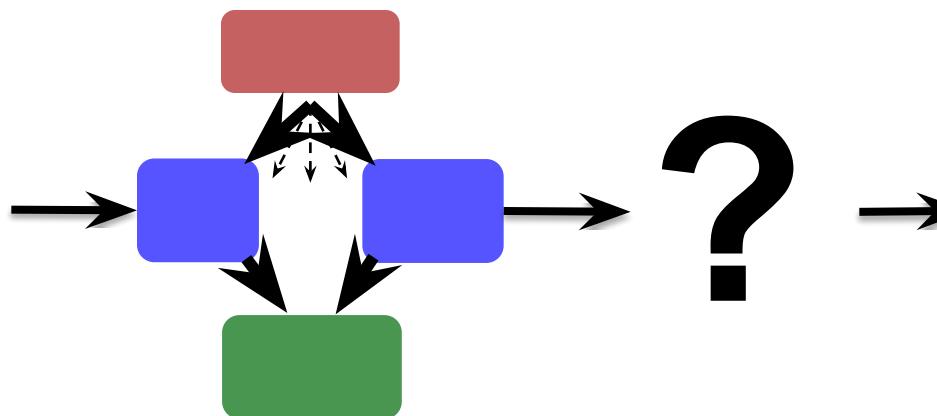
---



# End Up With This?

---

DAT  
A



# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. **Scale-up gradually**
5. Make it work consistently
6. What more can you automate or error-check?

(And remember to document!)

# Scaling Workflows

- Your (“small”) DAG runs! Now what?
  - Need to make it run *full scale*



to the  
moon!



# Scaling Up: Rules of Thumb

---

- CPU (single-threaded)
  - Best jobs run between **10 min** and **10 hrs**  
(Upper limit somewhat soft)
- Data (disk and network)
  - What is the balloon factor of your workflow?
  - Know when you are moving data and how to minimise it
- Memory
  - How much RAM / core do you have?

# Testing, Testing, 1-2-3 ...

---

- ALWAYS test a subset after making changes
  - How big of a change needs retesting?
- Scale up gradually
- Avoid making problems for others (and for yourself)

# Scaling Up - Things to Think About

---

- More jobs:
  - most submit queues will falter beyond ~10,000 total jobs
- Larger files:
  - more disk space, perhaps more memory
  - potentially more transfer and compute time

**Be kind to your submit and execute nodes  
and to fellow users!**

# Building a Good Workflow

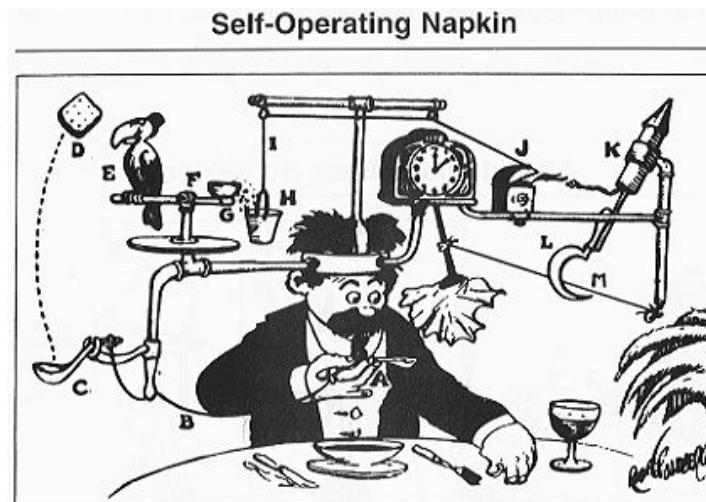
---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. **Make it work consistently**
6. What more can you automate or error-check?

(And remember to document!)

# Robust Workflows

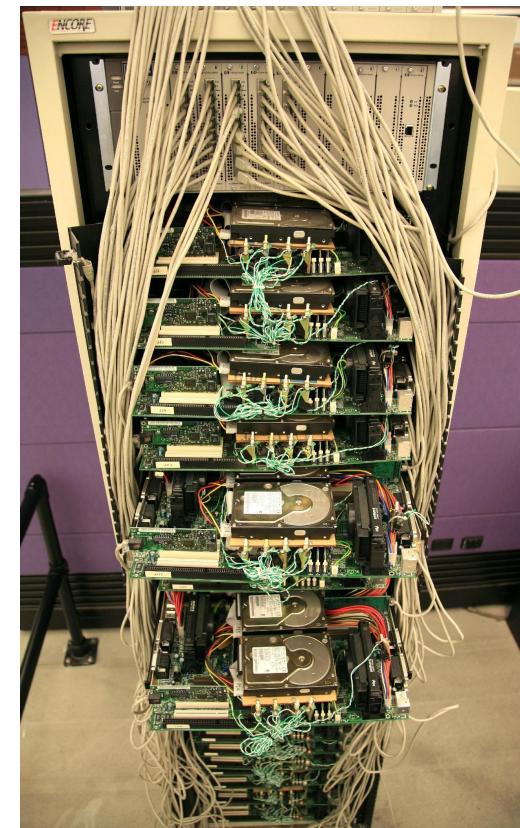
- Your workflow runs at scale! Now what?
  - Need to make it run *everywhere, everytime*
  - Need to make it run *unattended*
  - Need to make it run *when someone else tries*



# Make It Run Everywhere

---

- What does the computing environment have?
  - Prepare for very little
- Bring as much as possible with you, including:
  - the executable you need
  - the software it depends on
  - data dependencies??



# The Spectrum

---

- Laptop (1 machine)
  - You control everything!
- Local cluster (<100 - 1000 cores)
  - You can ask an admin nicely
- Cloud (core you pay for)
  - Be prepared to build your own cluster

# Make It Work Everytime

---

- What could possibly go wrong?
  - Eviction
  - Non-existent dependencies
  - File corruption
  - Performance surprises
    - Network
    - Disk
    - ...
  - *Maybe even a bug in your code*



# Make It Run(-able) for Someone Else

---

- Automation is a step towards making your research reproducible by someone else
  - Work hard to make this happen.
  - It's *their* throughput, too.
- Can benefit those who want to do similar work

# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. Make it work consistently
6. **What more can you automate or error-check?**

(And remember to document!)

# Automate All The Things?

---

- Well, not really, but kind of ...
- Really: What is the minimal number of manual steps necessary?  
even 1 might be too many; zero is perfect!
- Consider what you get out of automation
  - time savings (including less ‘babysitting’ time)
  - reliability and reproducibility

# Automation Trade-offs

HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE  
EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?  
(ACROSS FIVE YEARS)

		HOW OFTEN YOU DO THE TASK					
		50/DAY	5/DAY	DAILY	WEEKLY	MONTHLY	YEARLY
	1 SECOND	1 DAY	2 HOURS	30 MINUTES	4 MINUTES	1 MINUTE	5 SECONDS
	5 SECONDS	5 DAYS	12 HOURS	2 HOURS	21 MINUTES	5 MINUTES	25 SECONDS
	30 SECONDS	4 WEEKS	3 DAYS	12 HOURS	2 HOURS	30 MINUTES	2 MINUTES
HOW MUCH TIME YOU SHAVE OFF	1 MINUTE	8 WEEKS	6 DAYS	1 DAY	4 HOURS	1 HOUR	5 MINUTES
	5 MINUTES	9 MONTHS	4 WEEKS	6 DAYS	21 HOURS	5 HOURS	25 MINUTES
	30 MINUTES		6 MONTHS	5 WEEKS	5 DAYS	1 DAY	2 HOURS
	1 HOUR		10 MONTHS	2 MONTHS	10 DAYS	2 DAYS	5 HOURS
	6 HOURS				2 MONTHS	2 WEEKS	1 DAY
	1 DAY					8 WEEKS	5 DAYS



# Make It Work Unattended

---

- Remember the ultimate goal:  
**Automation! Time savings!**
- Potential things to automate:
  - Data collection
  - Data preparation and staging
  - Submission
  - Analysis and verification
  - Workflow testing



# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. Make it work consistently
6. What more can you automate or error-check?

**(And remember to document!)**

# Documentation at Multiple Levels

---

- In job files: comment lines
  - submit files, wrapper scripts, executables
- In README files
  - describe file purposes
  - define overall workflow, justifications
- In a document!
  - draw the workflow, explain the big picture

---

# PARTING THOUGHTS

# Getting Research Done

---

- End goal: getting the research done
- Hopefully you now have the tools to get the most out of:
  - **Computing**: which approach and set of resources suit your problem?
  - **High Throughput computing**: optimize throughput, use portable data and software
  - **Workflows**: test, automate and scale